

How Many Iterations are Enough? Bayesian Iterative Outlier Detection for Conservative Randomized Testing

Richard Carson

February 10, 2026

Abstract

Determining the number of iterations required for randomized tests to reliably detect edge cases is a common challenge in software testing. Traditional approaches—fixed iteration counts or fuzzing; either risk insufficient coverage or incur high computational costs, and are often unsuitable for continuous integration (CI) pipelines.

For numerical software, failures are frequently caused by extreme data points, and synthetic data with known outlier distributions is used to assess robustness. However, estimating the required number of test iterations in such scenarios remains difficult.

We introduce Bayesian Iterative Outlier Detection (BIOD), a method that combines sequential analysis with Bayesian updating to adaptively determine test duration based on observed failures and prior knowledge of outlier distributions. BIOD balances computational cost with confidence in test outcomes, is robust to poorly specified parameters, and produces reliable estimates of failure rates. Empirical evaluation shows that BIOD accurately models failures under independent, normally distributed errors, making it well suited for CI and other resource-constrained testing environments.

1 Introduction

1.1 Problem Statement

1.1.1 Mathematical Setup

Population and Sets

Let $P = \{X^{(1)}, X^{(2)}, \dots\}$ be a countably infinite population of sets. Each set is of fixed size n :

$$X^{(j)} = \{X_0^{(j)}, X_1^{(j)}, \dots, X_{n-1}^{(j)}\}, \quad j \in \mathbb{N}.$$

Each element $X_i^{(j)}$ is drawn i.i.d. and has probability p of being an **extreme outlier** (an event we denote $O_i^{(j)} = 1$ if $X_i^{(j)}$ is an outlier, 0 otherwise):

$$\Pr[O_i^{(j)} = 1] = p$$

Blackbox Function

There is a function

$$F : \mathcal{X} \rightarrow \{0, 1\}, \quad F(X^{(j)}) = \begin{cases} 1 & \text{if set passes} \\ 0 & \text{if set fails} \end{cases}$$

We assume there exists a relationship between the presence of outliers in $(X^{(j)})$ and failure, but the exact nature of this relationship is unknown:

$$\text{unknown mapping } F(X^{(j)}) \sim f(O_0^{(j)}, \dots, O_{n-1}^{(j)})$$

Sampling Problem

Let $(S \subseteq P)$ be the subset of sets actually observed/tested.

Define the event $E = \{\exists X^{(j)} \in S : F(X^{(j)}) = 0\}$, i.e., "we observe at least one failing set."

The goal is:

$$\text{Find the minimal } |S| \text{ such that } \Pr[E] \geq \gamma$$

where $(\gamma \in (0, 1))$ is a target confidence (e.g., 0.95).

1.1.2 Software Testing Perspective

When testing code consuming numerical data, it is common to use random transforms to produce synthetic data sets for testing.

However, this introduces the probability that any given run of the test will fail to catch edge cases, or fail due to a specific random seed.

Common approaches to mitigate this risk include fixed-budget testing, empirical convergence, fuzzing, or Monte Carlo estimation, all of which have their own drawbacks and limitations.

One major limitation of many approaches is the compute or time cost of running the test, making their use in CI pipelines or other time-sensitive contexts difficult.

2 Proposed Solution

2.1 Overview

To address the challenges outlined in Section 1, we introduce BIOD (Bayesian Iterative Outlier Detection), a method for efficiently estimating the minimal number of test iterations required to achieve a desired confidence level of detecting failure in code consuming numerical data, particularly while testing with synthetic data with known noise characteristics.

BIOD combines a modified Bayesian update procedure with a frequentist stopping criterion, and incorporates a damping factor to control conservatism in the initial estimate of failure probability. The method is highly tunable, has no systemic bias towards false positives or false negatives, and is robust to mis-specification of its input parameters, making it suitable for a wide range of use cases and risk tolerances.

The output of the test includes both a pass/fail result, and a lower bound confidence interval for the true failure rate among data sets, allowing users to make informed decisions about the reliability of their code and the potential risks associated with it. The method is designed to be computationally efficient, making it suitable for use in CI pipelines and other time-sensitive contexts.

2.2 Comparison to Other Methods

Below we outline existing approaches to the problem, and their limitations compared to the proposed BIOD method.

Fixed-budget testing

- *Approach* - Select a large fixed number of iterations k based on heuristics, and run the test for that many iterations.
- *Advantages* - This method is simple and easy to implement
- *Limitations* - It does not provide any guarantees about the confidence level of detecting failure, and can be computationally expensive if k is set too high. Additionally, in the

case of the absence of failure, this method provides no ability to stop early, and thus can waste resources.

Empirical convergence / burn-in heuristics

- *Approach* - Test until the observed pass ratio converges to a stable value, or until a certain number of consecutive passes are observed.
- *Advantages* - This method can be more efficient than fixed-budget testing, as it allows for early stopping if the test appears to be consistently passing.
- *Limitations* - It can be sensitive to the choice of convergence criteria, and may not provide guarantees about the confidence level of detecting failure. Additionally, very rare failures may not be observed within a reasonable number of iterations, leading to false confidence in the reliability of the code.

Worst-case / adversarial thinking

- *Approach* - Design tests to specifically target edge cases and potential failure modes, rather than relying on random sampling.
- *Advantages* - This can be effective in catching specific known issues, and can be more efficient than random sampling if the failure modes are well understood.
- *Limitations* - It may not provide a comprehensive assessment of the code's reliability, and can be time-consuming to design and implement. Additionally, it may not be effective in catching unknown failure modes or edge cases that were not anticipated during test design.

Classical statistical power analysis

- *Approach* - Calculate the required sample size to achieve a desired power level for detecting a specific effect size, based on assumptions about the underlying distribution of the data and the expected failure rate.
- *Advantages* - This can provide a more rigorous basis for determining the number of iterations required, and can be tailored to specific tests and failure modes.
- *Limitations* - It relies heavily on the accuracy of the assumptions made, and may not be robust to mis-specification of parameters. Additionally, it may not be practical in cases where the underlying distribution is unknown or difficult to model and is specific to the particular test being run, making it less generalizable.

Monte Carlo estimation

- *Approach* - Run a large number of random simulations to estimate the probability of failure, and using this estimate to determine the required number of iterations to achieve a desired confidence level.
- *Advantages* - This can be effective in providing a more accurate estimate of the failure probability, and can be tailored to specific tests and failure modes.
- *Limitations* - It can be computationally expensive, particularly for complex tests or large data sets. Additionally, it may not provide guarantees about the confidence level of detecting failure, and can be sensitive to the choice of simulation parameters.

Fuzzing

- *Approach* - Generate random or semi-random inputs to test the code, with the goal of uncovering edge cases and potential failure modes.
- *Advantages* - Very comprehensive; can be effective in catching a wide range of issues, including those that were not anticipated during test design.
- *Limitations* - Computationally expensive, as it requires a large number of test iterations to achieve a high confidence level of detecting failure. Additionally, it may not be suitable for use in CI pipelines or other time-sensitive contexts, as it can lead to long test execution times.

2.3 Algorithm Details

The BIOD method consists of three main phases, depicted in Figure 1: the estimation phase, the iterative testing phase, and the output phase.

We begin by defining the following parameters and constants used throughout the algorithm:

- n The size of the data set under test.
- p The probability of any individual point being an extreme outlier, chosen conservatively (underestimation of the true risk).
- q The desired confidence that enough iterations have been performed to detect a theoretical bug caused by an extreme outlier.
- `pass_ratio` The required ratio of passes in the observed sample of iterations to consider the test a success.
- $C = 1.864$ A damping constant reflecting the diminishing impact of additional data points on the likelihood of failure due to extreme outliers




Figure 1: Flowchart of the BIOD testing procedure.

$p_s = 3.35 \cdot 10^6$ A strength parameter for the Bayesian prior, reflecting the confidence in the initial estimate of the failure probability

Derivations for suggested values constant are included in Section 2.5.

q represents the target confidence γ of observing a failing set within the sample population ($\Pr[E] \geq q$). It should be selected to be a reasonably conservative value, as defined in Section 2.5.

p should be selected to be a very conservative underestimate of the true point-wise probability of extreme outliers, to ensure that the algorithm does not underestimate the required number of iterations.

q should be selected based on the user's tolerance for risk, with higher values indicating a lower tolerance for risk. For many users, `pass_ratio` and q could be conflated, since they both reflect differing aspects of the user's tolerance for risk. However they are kept separate to allow for more advanced use cases where the user may wish to set a high confidence level for catching edge cases, but a lower pass ratio to allow for known issues or an allowable non-zero failure rate. The advantage of this is that the user can be more certain that any failures observed are due to real issues, rather than random chance.

The sections below describe the 3-phase approach for the BIOD method.

2.3.1 Estimation Phase

Given n the number of items per set X , and p the point-wise lower-bound estimate for outlier frequency, we calculate an initial conservative estimate for the set-wise probability of failure p_{fail} :

$$p' = \frac{p}{(1 + C \cdot n \cdot p)}, p_{\text{fail}} = 1 - (1 - p')^n \quad (1)$$

p' is used to model the saturation of the effect of individual outliers on set-wise failure. It represents the inverse relationship $\frac{1}{np}$ between number of expected outliers, and the effects of iterations between them. For more on the derivation of this form, and the recommended value for $C = 1.864$, see Section 2.5.1. We call the the ratio applied to p the "damping fraction", since it serves to dampen the effect of increasing n on p_{fail} .

For small n the probability of failure is roughly linear in n , consistent with a binomial model where $p_{\text{fail}} \approx 1 - (1 - p)^n$. As n becomes large, the damping fraction saturates, and the effective failure probability approaches a limiting value $p_{\text{fail}} = 1 - e^{-1/C}$. In this regime, the number of failures per iteration can be approximated by a Poisson distribution with $\lambda = \frac{1}{C}$, reflecting the diminishing returns of increasing data set size on the likelihood of failure $F(X)$ due to the presence of extreme outliers in the set.

An initial estimate for the number of iterations k required to achieve the desired confidence q of observing at least one failing set in the collected sample S (that is to say the likelihood estimate for $\Pr[E]$ given q) can then be computed as:

$$k = \left\lceil \frac{\ln(1 - q)}{\ln(1 - p_{\text{fail}})} \right\rceil \quad (2)$$

Where p_{fail} is the probability of any single data set failing the test.

Finally, Bayesian priors for p_{fail} are calculated using p_s :

$$\alpha = p_{\text{fail}} \cdot p_s, \quad \beta = (1 - p_{\text{fail}}) \cdot p_s \quad (3)$$

For more on the derivation of this form, and the recommended value for $p_s = 3.35 \cdot 10^6$, see Section 2.5.2.

2.3.2 Iterative Testing Phase

In this phase we will track the following variables:

k The current estimate for iterations required to achieve the given confidence

α, β The priors for our Bayesian updates

iterations The number of iterations performed so far

passes The number of successful iterations

passes_since_failure A tally of successes since the last failure

The test involves running a blackbox testing routine on a new set of size n , $X^i \in P$, in each iteration in order to evaluate $F(X^i)$. We recommend each test utilize a seeded transform on an initial input set.

For each iteration, perform the following steps:

Stopping Criteria

The current value for k represents the number of iterations required to have confidence q of having observed at least one failure, given the current estimate for set-wise failure probability p_{fail} .

Thus, if $iterations \geq k$, we can be confident that we have observed at least one failure if the true failure probability is equal to or greater than our current estimate, and thus we can stop the test and move on to the output phase.

Timeout Criteria

To prevent infinite loops in the case of extremely low failure probabilities, we recommend implementing a timeout mechanism. If the total elapsed time for the testing process exceeds a predetermined threshold, the test should terminate and report failure due to insufficient evidence to support a claim.

Update Procedure

Note:

With our assumptions that p_{fail} represents a conservative underestimate, and that at least one failing set exists, we can attribute any observed successes prior to a failure to chance, and therefore the update procedure should only be performed on failure, since this is the only time we can be sure that we have observed a real failure rather than just random chance. Additionally, this prevents k from growing in the base-case where no failures exist.

- Increment *iterations* by 1 to reflect the new test iteration, and perform the blackbox test on the new set. The test should determine if the set passes or fails.
- If the set passes, we increment *passes* and *passes_since_failure* by 1, and repeat the procedure for the next iteration.
- If the set fails, we update our estimates:
 - Increment *alpha* by 1, to represent the detected failure
 - Perform $\beta+ = passes_since_failure$, and reset $passes_since_failure = 0$
 - Recalculate the estimated set-wise failure risk: $p_{\text{fail}} = \frac{\alpha}{\alpha+\beta}$

- Update the number of iterations required: $k = \left\lceil \frac{\ln(1-q)}{\ln(1-p_{\text{fail}})} \right\rceil$
- Repeat the iteration procedure

2.3.3 Output Phase

Once the required evidence has been collected, we can calculate the observed ratio of passes to failures r_{obs} as:

$$r_{\text{obs}} = \frac{\text{total passes}}{\text{total iterations}} \quad (4)$$

The final test result can then be determined by comparing r_{obs} to the user-defined `pass_ratio`:

$$\text{test result} = \begin{cases} \text{pass} & \text{if } r_{\text{obs}} \geq \text{pass_ratio} \\ \text{fail} & \text{if } r_{\text{obs}} < \text{pass_ratio} \end{cases}$$

To return a confidence interval for the true failure rate among data sets, we can construct a new beta distribution using the observed data, with parameters $\alpha' = \text{total failures} + 1$ and $\beta' = \text{total passes} + 1$.

$$p_{\text{fail,lower}} = F^{-1} \left(\frac{1-q}{2}; \text{total failures} + 1, \text{total passes} + 1 \right) \quad (5)$$

$$(6)$$

where F^{-1} is the Beta quantile function.

We do not use the beta distribution from the iterative phase, since this distribution is heavily influenced by the initial conservative estimate for p_{fail} , and thus will not reflect the observed pass ratio. Instead, we use a new beta distribution based solely on the observed data to calculate a credible interval for p_{fail} . Because the goal is to discover the number of iterations in which we expect exactly 1 failure, only the lower bound of the credible interval is relevant

The outputs for the test are then:

- test result The test result
- r_{obs} , the observed pass ratio
- The calculated lower bound on $p_{\text{fail, true}}$ given q

These can be communicated succinctly as: “[$100 \cdot r_{\text{obs}}$]% of tests passed. [$100 \cdot q$]% confidence that the true failure rate among datasets is at least $p_{\text{fail,lower}}$.”

2.4 Limitations

The proposed method has several limitations and assumptions worth noting:

- The size of the data set n must be known in advance, and should remain constant throughout the testing process.
- The method assumes that the probability of extreme outliers is independent and normally distributed across the data set, which may not hold true in all cases.
- Very small n (< 100) require a very large number of iterations to converge; this is an expected bi-product of the model used, and we therefore recommend $n \geq 100$

2.5 Derivations

2.5.1 The Damping Constant - C

As n increases, the naive binomial form $p_{\text{fail}} = 1 - (1 - p')^n$ for expected failure trends to 100%. To resolve this, we introduce the damping fraction $p' = \frac{p}{(1+C \cdot n \cdot p)}$

We began with the form

$$p' = \frac{p}{n \cdot p} \quad (7)$$

In order to eliminate the undefined case at $n = 0$, and to ensure that $p' > p$ in all cases, we modified this to the form:

$$p' = \frac{p}{(1 + n \cdot p)} \quad (8)$$

And finally introduced the constant C to allow for tuning of the saturation effect:

$$p' = \frac{p}{(1 + C \cdot n \cdot p)} \quad (9)$$

C represents the binding strength modulating between outliers, and thus the rate of saturation of the effect of increasing n on p_{fail} .

The effect is the introduction of an asymptotic floor on the minimum number of iterations required for a set as $n \rightarrow \infty$, k_{floor} .

To find a reasonable value for C , we define the reasonable universe of parameters for the test space. For a theoretical set of infinite size, there is naturally a minimum number of iterations under which the test is not useful - k_{floor} . There is also a practical maximum number, above which the computation required for extremely large sets becomes infeasible.

We will define that practical range for k_{floor} to be between 3 and 10 iterations, and the we will define the lowest confidence q we consider reasonably conservative to be 0.8.

We we can analyze the effect of C effect on k_{floor} as n approaches infinity:

$$p' = \frac{p}{1 + Cnp} \quad (10)$$

$$p_{\text{fail}} = 1 - (1 - p')^n \quad (11)$$

For large n , $(1 - p')^n \rightarrow e^{-np'}$, so

$$p_{\text{fail}} = 1 - e^{-\frac{np}{1+Cnp}} = 1 - e^{-\frac{1}{C} \cdot \frac{Cnp}{1+Cnp}} \quad (12)$$

As $n \rightarrow \infty$, $\frac{Cnp}{1+Cnp} \rightarrow 1$, giving us $p_{\text{fail}} = 1 - e^{-1/C}$

Then

$$k = \frac{\ln(1 - q)}{\ln(1 - p_{\text{fail}})} \Rightarrow k_{\text{floor}} = \frac{\ln(1 - q)}{\ln(e^{-1/C})} = \frac{\ln(1 - q)}{-1/C} = -C \ln(1 - q) \quad (13)$$

Solving for C :

$$C = \frac{k_{\text{floor}}}{-\ln(1 - q)} \quad (14)$$

If we set $k_{\text{floor}} = 3$ as a reasonable minimum number of iterations that could be considered useful for a lower-bound on the value for q we consider reasonably conservative of $q = 0.8$, we find:

$$C = \frac{k_{\text{floor}}}{-\ln(1 - q)} \Rightarrow C = \frac{3}{-\ln(1 - 0.8)} \approx 1.864$$

Additionally, since $k_{\text{floor}} = -C \cdot \ln(1 - q)$, we can measure the effect of varying C on k_{floor} for a given q :

$$k_{\text{floor}} = -C \ln(1 - q)$$

$$k_{\text{floor}} = -(1.864 - 0.5) \ln(1 - 0.8) \approx 2.19527$$

$$k_{\text{floor}} = -(1.864) \ln(1 - 0.8) \approx 2.99999$$

$$k_{\text{floor}} = -(1.864 + 0.5) \ln(1 - 0.8) \approx 3.80471$$

Taking into account the effect of ceil in the true calculation of k_{initial} , we can see that variations in C of approximately ± 0.5 around the nominal value of 1.864 produces a difference of approximately ± 1 in the final value, showing that this method is reasonably robust to small variations in C .

We will therefore recommend $C = 1.864$ as the nominal value for the damping constant.

To verify this value, we generated the k_{floor} produced by a variety of values within the defined parameters space, and then took C_{mean} and C_{σ} to produce the table and graph shown in Figure 5.

From the table we can see that the $C \approx 1.8$ produces values right in our defined practical range for k_{floor} across a variety of confidence levels, confirming our derived value.

If a use-case requires a different practical range for k_{floor} or q , a value for C should be selected using the formula above (13).

2.5.2 Prior Strength - ps

Isolating the effects of parameters on bias

Our first step to determining a reasonable p_s for a given model was to run the experiment described in Section 3.5, varying p_{fail} and pass_ratio over the base-case of $p_s = 1$, to determine if either parameter had a significant effect on bias.

These parameters were chosen since they are calculated directly from n and p , and thus would allow us to infer the effect of those parameters on bias.

For this and all experiments in this section we define bias as the difference between the test’s observed pass ratio r_{obs} , and the true underlying pass rate for the simulated trial, r_{true}

$$\text{bias} = r_{\text{obs}} - r_{\text{true}} \quad (15)$$

A positive bias indicates the test overestimated the pass rate (false negative), while a negative bias indicates the test underestimated the pass rate (false positive).

Our hypothesis was that the delta between the expected level of failure p_{fail} and the acceptable level of failure $1 - r_{\text{pass}}$ would be a good predictor of bias.

However, as can be seen in Figure 6, neither parameter had a measurable effect on bias - which appeared to be randomly distributed around 0 with a constant error regardless of the values of p_{fail} and pass_ratio . Since p_{fail} encodes for k , and contains both n and p , we can infer that none of n , p , p_{fail} , or k have a significant effect on bias, which appears to be a set function for any inputs into the model given an unknown underlying failure rate.

Effect of prior strength on bias and compute

Our next step was to run a simplified experiment varying only p_s , shown in Figure 8a (a), to determine the effect of p_s on bias. This demonstrated the bias was indeed centered at 0, with a variance that decreased with increasing p_s , plateauing near $p_s = 1e7$.

We next ran the same experiment measuring the fraction of initially estimated iterations required to converge at various values of p_s , shown in Figure 8b (b). This demonstrated that increasing p_s also increased the required computation, again reaching an asymptotic limit, this time around $p_s = 1e8$.

Modeling the tradeoff between bias and compute

Following Experiment 3, we fit the relationship between the performed fraction of maximum iterations and p_s to a pair of Hill equations [2] of the form:

$$y = y_{min} + (y_{max} - y_{min}) \frac{x^n}{k^n + x^n} \quad (16)$$

Where:

- y is the observed variable (bias stddev or compute fraction)
- x is p_s , the prior strength
- n is the Hill coefficient, which determines the steepness of the curve
- k is the value of x at which y reaches half of its maximum value
- y_{min} and y_{max} are the starting and ending values of y as x increases

Bias was first normalized to the maximum observed value, and compute was already a fraction of the maximum, making both variables suitable for fitting to this form.

The hill curve was chosen due to the need for a monotonic, saturating sigmoid function.

This resulted in a good fit for both bias and compute ($R^2 > 0.998$ for both curves).

The fits had the following parameters:

- Bias
 - $n = 0.548$
 - $k = 4.77 \cdot 10^4$
 - $y_{min} = 1.0$
 - $y_{max} = 0.323$
- Compute
 - $n = 0.639$
 - $k = 6.70 \cdot 10^5$
 - $y_{min} = 0.111$
 - $y_{max} = 1.0$

Selecting a recommended prior strength

We then defined a Lagrangian objective function [1] to balance the competing goals of minimizing bias and minimizing compute, introducing β , the tradeoff coefficient between the two objectives representing the relative importance of compute savings versus bias reduction:

$$\text{objective_score}(p_s) = \text{bias_std}(p_s) - \beta \cdot \text{compute_frac}(p_s) \quad (17)$$

By sweeping β over logarithmic space from 10^{-5} to 10^1 , or in other words, from a strong preference for minimizing bias to an equal weighting of both objectives, and solving for the root of the objective function, we can build a table of optimal p_s values for various tradeoff preferences.

The result, shown in Figure 10, demonstrates a flat region of optimal p_s values centered at $3.35 \cdot 10^6$, in which bias becomes relatively insensitive to p_s , while compute savings continue to improve significantly as p_s decreases. The tradeoff curve and suggested optimal value are shown in Figure 9.

We therefore recommend a nominal value of $p_s = 3.35 \cdot 10^6 \pm 1.98 \cdot 10^4$ for a good balance between compute and bias. Compared to the largest reasonable value of $p_s = 1e7$, this value results in only a 3% increase in bias stddev, while saving 24% of worst-case compute. By comparison, the less conservative value of $p_s = 4.43e5$ results in a 12% increase in bias stddev, while saving 39% of worst-case compute.

The recommendation corresponds to $\beta = 0.51$, meaning that compute savings are considered roughly half as important as bias reduction.

It is noteworthy that since compute is a fraction of the initially estimated iterations actually completed, this will already scale with n ; The larger the set size, the greater the potential savings as compared to the initially estimated worst-case compute. Thus, this recommended value for p_s should be robust across a wide range of n values.

3 Methodology

3.1 Experiment 1: Conservative Estimate of p_{fail} vs Simulated True p_{fail}

Objective. To verify that calculated p_{fail} is a conservative underestimate of the true simulated p_{fail} across a range of dataset sizes n .

3.1.1 Setup.

- Vary set size n logarithmically from $1e0$ to $1e5$.
- For each n , run 10,000 trials.
- For each trial, simulate failure due to extreme outliers, accounting for interactions between outliers:
 - Draw n samples from a standard normal distribution ($\mu = 0, \sigma = 1$).




Figure 2: Graph showing the conservative estimates of p_{fail} vs simulated true p_{fail} for various dataset sizes n

- Define extreme outliers as samples with absolute value greater than 3.
- A trial fails if the sum of all extreme outliers exceeds 3 in absolute value.
- Record the observed failure rate across trials as the simulated true p_{fail} .
- For each n , calculate the corresponding conservative initial estimate of p_{fail} using the BIOD method:
 - Confidence level: 0.99
 - Outlier probability: 0.0027 (approximate probability of $|x| > 3\sigma$ in a normal distribution)
 - Pass ratio: 0.0 (not used in this experiment)

3.1.2 Definition of Metrics.

- $p_{\text{fail, true}}$: Observed failure rate from simulations.
- $p_{\text{fail, calc}}$: Calculated conservative estimate of p_{fail} using BIOD.

3.1.3 Hypothesis.

BIOD is expected to yield a calculated $p_{\text{fail, calc}}$ that is less than or equal to the observed $p_{\text{fail, true}}$ across all tested set sizes n .

3.1.4 Results.

As shown in Figure 2, $p_{\text{fail, calc}}$ consistently underestimates $p_{\text{fail, true}}$ across the full range of dataset sizes n . Moreover, the two curves exhibit similar functional form, indicating that the estimator preserves the underlying dependence on n while remaining systematically conservative.

3.1.5 Analysis / Interpretation.

The results confirm the hypothesis that the BIOD-calculated p_{fail} is a conservative underestimate of the true simulated p_{fail} across all dataset sizes tested.

By following the same trend as the true p_{fail} , the BIOD estimate avoids the trap of underestimating by enough to make initial k too large to be useful, as n increases, while still providing a reliable conservative estimate.

3.1.6 Summary / Recommendation.

The BIOD method provides a reliable conservative estimate of p_{fail} , without trend inversion as n varies, making it suitable for this application.

3.2 Experiment 2: Predicted k to Observed k to Failure




Figure 3: Graph showing the predicted number of iterations k to observed failure vs the simulated results

Objective. To verify that the predicted number of iterations k to failure from BIOD encompasses the observed number of iterations k to failure across a range of dataset sizes n .

3.2.1 Setup.

- Vary set size n logarithmically from $1e0$ to $1e5$.
- For each n , run 10,000 trials.
 - Draw n samples from a standard normal distribution ($\mu = 0, \sigma = 1$).
 - Define extreme outliers as samples with absolute value greater than 3.
 - A trial fails if the sum of all extreme outliers exceeds 3 in absolute value.
- For each trial, record the number of iterations k taken to find a failing set
- Take the mean and standard deviation of observed k across trials.
- After all trials for a given n , calculate the predicted k to failure using the BIOD method:
 - Confidence level: 0.99
 - Outlier probability: 0.0027 (approximate probability of $|x| > 3\sigma$ in a normal distribution)
 - Pass ratio: 0.0 (not used in this experiment)

3.2.2 Definition of Metrics.

- $k_{\text{obs, mean}}$: Mean observed iterations to failure from simulations.
- $k_{\text{obs, stdev}}$: Standard deviation of observed iterations to failure from simulations.
- $k_{\text{predicted}}$: Predicted iterations to failure using BIOD.

3.2.3 Hypothesis.

The BIOD is expected to yield a predicted $k_{\text{predicted}}$ that is greater than or equal to the mean observed $k_{\text{obs, mean}}$ across all tested set sizes n .

3.2.4 Results.

As shown in Figure 3, $k_{\text{predicted}}$ consistently overestimates $k_{\text{obs, mean}}$ across the full range of dataset sizes n . The predicted values remain above the observed means, following the same trend as observed k to failure, similarly to Experiment 1.

The observed standard deviation $k_{\text{obs, stdev}}$ shrinks as n increases, indicating more consistent behavior in larger datasets.

3.2.5 Analysis / Interpretation.

The results confirm the hypothesis that the BIOD-predicted k to failure is a conservative overestimate of the mean observed k to failure across all dataset sizes tested.

This means that for a well specified test, BIOD’s predicted k will likely include a failure within that number of iterations, providing confidence in the method’s reliability.

3.2.6 Summary / Recommendation.

The BIOD method provides a reliable conservative prediction of iterations to failure, making it suitable for this application.

3.3 Experiment 3: Confidence Interval on p_{fail} Estimates




Figure 4: Graph showing the mean and standard deviation of lower-bound confidence intervals on the sample of iterations observed by BIOD from a simulation of 10,000 trials, compared to the simulated true p_{fail} for various dataset sizes n , and for various values of point-wise p . Green was deliberately underestimated, blue was exact, and red was deliberately overestimated. Markers indicate the number of iterations k performed by BIOD at each point.

Objective. To verify that the calculated p_{fail} from BIOD provides accurate confidence intervals that encompass the observed p_{fail} across a range of dataset sizes n and point-wise p values.

3.3.1 Setup.

- Vary set size n logarithmically from $1e0$ to $1e5$.
- Test 10,000 sets for each combination of n , to calculate observed p_{fail} .
- Run BIOD with the same results to calculate confidence intervals on p_{fail} for each value of point-wise p :
 - Underestimate of p : $p_{\text{outlier}} = 1e - 4$
 - Exact p : $p_{\text{outlier}} = 0.0027$ (approximate probability of $|x| > 3\sigma$ in a normal distribution)
 - Overestimate of p : $p_{\text{outlier}} = 0.01$

3.3.2 Definition of Metrics.

- p_{fail} : True failure rate from simulations.
- $p_{\text{fail, biod, min}}$: Lower bound of BIOD-calculated confidence interval for p_{fail} .
- $p_{\text{fail, biod, max}}$: Upper bound of BIOD-calculated confidence interval for p_{fail} .
- $p_{\text{fail, obs}}$: Observed failure rate in the sample BIOD selects.

3.3.3 Hypothesis.

The BIOD-calculated confidence intervals for p_{fail} are expected to encompass the observed p_{fail} for accurate and underestimated point-wise p values, and to potentially miss for overestimated point-wise p values, across all tested set sizes n .

3.3.4 Results.

As shown in Figure 4, the calculated confidence intervals for p_{fail} from BIOD successfully encompass the observed p_{fail} for both all tested set sizes n even for badly specified point-wise p values.

When point-wise p is underestimated (green), the confidence intervals are narrower and more accurate, closely tracking the observed p_{fail} .

When point-wise p is exact (blue), the confidence intervals remain accurate, but is much wider than the underestimated case.

When point-wise p is overestimated (red), the confidence intervals are significantly wider for small n and converge to the exact case as n increases, but still encompass the observed p_{fail} .

3.3.5 Analysis / Interpretation.

These results show that BIOD is robust to badly specified point-wise p parameters, but gains significant accuracy when point-wise p is well specified or conservative (i.e., underestimated).


An overly conservative (underestimated) point-wise p leads to tighter confidence intervals, due to the larger number of iterations BIOD performs in this case, improving the accuracy of the p_{fail} estimate.

An overly aggressive (overestimated) point-wise p leads to wider confidence intervals, especially for small n , due to the smaller number of iterations BIOD performs in this case, reducing the accuracy of the p_{fail} estimate. However, as n increases, the effect of point-wise p mis-specification diminishes, and the confidence intervals converge towards those obtained with exact point-wise p .

3.3.6 Summary / Recommendation.

The BIOD method provides reliable confidence intervals for p_{fail} estimates, particularly when point-wise p is well specified or underestimated, but is still useful when point-wise p is overestimated.

3.4 Experiment 4: Damping Constant Effect on k floor



$C \backslash q$	0.8	0.9	0.95	0.99
1.56	3	4	5	8
1.82	3	5	6	9
2.07	4	5	7	10
2.33	4	6	7	11
2.59	5	6	8	12
2.84	5	7	9	14
3.1	5	8	10	15
3.36	6	8	11	16
3.61	6	9	11	17
3.87	7	9	12	18
4.13	7	10	13	20

Figure 5: Left: Graph of k_{floor} vs C values in the defined parameter space. Right: Table showing values of k_{floor} for $C_{\text{mean}} \pm C_{\sigma}$ for various confidence levels q .

Objective. To investigate the relationship between the damping constant C , desired k floor and confidence level q .

3.4.1 Setup.

The experiment will be performed in 2 parts:

Part 1 Vary q smoothly from 0.8 to 0.99 in increments of 0.01. For each q , vary k from 3 to 10 in increments of 1. Compute the corresponding C value for each (k, q) pair using the formula provided below:

$$C = \frac{k}{-\ln(1 - q)}$$

This will provide a range of C values corresponding to the desired k floor and q values defined in the parameter space.

Part 2 Use the results from Part 1 to compute the mean C_{mean} and standard deviation C_{σ} of the damping constant C across all (k, q) pairs. For the range $C = C_{\text{mean}} \pm C_{\sigma}$, compute the corresponding k_{floor} values for specific confidence levels $q = [0.8, 0.9, 0.95, 0.99]$:

$$k_{\text{floor}} = C \cdot -\ln(1 - q)$$

This will provide a table of k_{floor} values for the defined range of C and specific confidence levels.

3.4.2 Definition of Metrics.

Damping constant C for each desired k floor and pass probability q .

3.4.3 Hypothesis.

We expect C to increase monotonically with k and q .

3.4.4 Results.

The results, shown in Figure 5, confirm that C increases with both desired k floor and pass probability q , ranging from < 1 to > 6 in our defined parameter space.

Narrowing the range to $C = C_{\text{mean}} \pm C_{\sigma}$ gives us a range of 1.56 to 4.13, where k_{floor} values between 3 and 10 correspond to C values between 1.82 and 2.07 for confidence levels between 0.8 and 0.99.

3.4.5 Analysis / Interpretation.

The results confirm the hypothesis that C increases with both desired k floor and pass probability, and gives us a practical basis to verify the value found analytically in Section 2.5.1.

3.4.6 Summary / Recommendation.

From the table, a value between $C = 1.82$ and $C = 2.07$ best represents the desired k floor range of 3-10 for confidence levels between 0.8 and 0.99.

3.5 Experiment 5: Bias vs p_{fail} and ratio_pass




Figure 6: Graph demonstrating the lack of correlation between p_{fail} , pass_ratio , and bias in the base case of $p_s = 1$

Objective. To investigate the relationship between bias, p_{fail} , and pass ratio. p_{fail} is calculated using p , q and n , and so between it and pass_ratio we can explore the bias behavior of BIOD across the parameter space in the base-case of $p_s = 1$.

3.5.1 Setup.

- We select $\text{pass_ratio} \in \{0.2, 0.5, 0.8, 0.9, 0.95, 0.99, 0.999\}$.
- We vary p_{fail} from 0.05 to 0.50 in increments of 0.01
- For each p_{fail} , we compute the corresponding dataset size n using:

$$n = \left\lceil \frac{-\ln(1 - p_{\text{fail}})}{p'} \right\rceil, \quad p' = \frac{1}{1 + C \cdot n \cdot p}, \quad C = 1.864, \quad p = 10^{-6}$$

For each $(p_{\text{fail}}, \text{ratio_pass})$ pair, perform the 10,000 iterations of the following:




Figure 7: Scatter plot showing the distribution of true p_{fail} values across all iterations of the experiment, demonstrating the coverage of the parameter space.

- Draw a value x from a uniform distribution $U(0, 1)$.
- Determine range for the rate of failure based on the following piecewise function:

$$\text{failure_rate} = \begin{cases} [0.001, 0.01] & x < 0.1 \\ [0.01, 0.05] & 0.1 \leq x < 0.3 \\ [0.05, 0.20] & 0.3 \leq x < 0.65 \\ [0.20, 0.50] & 0.65 \leq x < 0.9 \\ [0.50, 0.999] & x \geq 0.9 \end{cases}$$

- Draw a true failure probability $p_{\text{fail, true}}$ from a uniform distribution over the determined range.
- Run BIOD using the failure rate found, noting the `pass_ratio` value at the end of the run, the initial estimate for k and the number of iterations actually performed.
- Compute bias as $p_{\text{fail, true}} - \text{pass_ratio}$.
- Compute the mean and standard deviation of bias across all iterations for each `pass_ratio` value.

For each p_{fail} , compute:

- The mean of all bias means across all `pass_ratio` values.
- The standard deviation of all bias means across all `pass_ratio` values.

- The mean of all bias standard deviations across all `pass_ratio` values.
- The standard deviation of all bias standard deviations across all `pass_ratio` values.

3.5.2 Definition of Metrics.

- $p_{\text{fail, true}}$: True failure probability drawn from the defined piecewise distribution, magnifying the presence of segments of the parameter space that are more likely to be encountered in practice, while still covering the full range of p_{fail} values.
- p_{fail} : Calculated failure probability from BIOD.
- `pass_ratio`: Pass ratio parameter provided to BIOD.
- *bias*: Difference between true failure probability and observed pass ratio. The mean of bias represents systematic error, while the standard deviation represents variability in the error.

3.5.3 Hypothesis.

We expect bias to increase proportionally with the difference between p_{fail} and `pass_ratio`.

If p_{fail} has a statistically significant effect on bias, we would expect to see trends in bias as p_{fail} increases, either in the mean, or in the standard deviation, or both.

If `pass_ratio` has a statistically significant effect on bias, we would expect to see large error bars on the mean or standard deviation of bias within each set of trials for each `pass_ratio` value.

3.5.4 Results.

Contrary to our hypothesis, we find that bias remains relatively constant across the parameter space, with mean bias values clustering around 0.0 and standard deviation of bias remaining low for all combinations of p_{fail} and `pass_ratio`.

We also found that the standard deviation of bias mean and standard deviation of bias standard deviation across all `pass_ratio` values remained low for all p_{fail} values.

3.5.5 Analysis / Interpretation.

Since the parameters tested involve all of BIOD's inputs except for prior strength, this experiment establishes that BIOD's bias behavior is not predictable from inputs, and will be a factor of prior strength alone.

The lack of correlation between bias and p_{fail} implies a lack of a statistically significant effect of p_{fail} on bias, and the mean near 0.0 across all `pass_ratio` values implies no systemic bias in either direction across the parameter space.

The small standard deviation of bias mean and bias standard deviation across all pass_ratio values implies no significant effect of pass_ratio on bias, nor any statistically significant interaction effects between pass_ratio and p_{fail} on bias.

3.5.6 Summary / Recommendation.

These results indicate that prior strength is the key parameter influencing bias in BIOD, and should be the focus of further investigation.

It suggests that p_s is not a function of inputs, but a constant that can be pre-computed to achieve desired bias characteristics.

3.6 Experiment 6: Prior Strength vs Bias and Compute




Figure 8: Comparison of bias behavior and convergence across p_s values

Objective. To investigate the relationship between prior strength, bias, and computational cost in BIOD.

3.6.1 Setup.

- We vary prior strength p_s from $10e0$ to $10e12$ in logarithmic increments of $10e0.5$.
- For each p_s , we perform 10,000 iterations of the following:
 - Draw a value x from a uniform distribution $U(0, 1)$.

- Determine range for the rate of failure based on the following piecewise function:

$$\text{failure_rate} = \begin{cases} [0.001, 0.01] & x < 0.1 \\ [0.01, 0.05] & 0.1 \leq x < 0.3 \\ [0.05, 0.20] & 0.3 \leq x < 0.65 \\ [0.20, 0.50] & 0.65 \leq x < 0.9 \\ [0.50, 0.999] & x \geq 0.9 \end{cases}$$

- Draw a true failure probability $p_{\text{fail, true}}$ from a uniform distribution over the determined range.
- Run BIOD using the failure rate found, noting the `pass_ratio` value at the end of the run, the initial estimate for k and the number of iterations actually performed.
- Compute bias as $p_{\text{fail, true}} - \text{pass_ratio}$.
- Compute the fraction of initial k actually performed: $\frac{k_{\text{performed}}}{k_{\text{initial}}}$.
- Compute the mean and standard deviation of bias across all iterations.
- Compute the mean and standard deviation of fraction of initial k performed across all iterations.
- Repeat the above procedure 50 times, to get the mean and standard deviation of the above metrics across runs for each p_s value.

3.6.2 Definition of Metrics.

- $p_{\text{fail, true}}$: True failure probability drawn from the defined piecewise distribution, magnifying the presence of segments of the parameter space that are more likely to be encountered in practice, while still covering the full range of p_{fail} values.
- p_s : Prior strength parameter provided to BIOD.
- *bias*: Difference between true failure probability and observed pass ratio. The mean of bias represents systematic error, while the standard deviation represents variability in the error.
- *compute_fraction*: Fraction of initial k actually performed, representing computational cost.

3.6.3 Hypothesis.

We expect bias to remain centered around 0.0, but its standard deviation to decrease as prior strength p_s increases. We also expect computational cost to increase with prior strength, as higher p_s values lead to larger initial k estimates.

3.6.4 Results.

As shown in Figure 8b, the mean bias remains centered around 0.0 across all prior strength values tested, confirming part of our hypothesis.

However, while the standard deviation of bias does decrease with increasing prior strength, there is a ceiling effect observed around $p_s = 10e7$ beyond which further increases in prior strength yielded no additional reduction in bias variability, while still increasing computational cost.

3.6.5 Analysis / Interpretation.

There appears to be a sigmoidal saturation effect in bias standard deviation with respect to both bias standard deviation and fraction of worst-case compute performed.

This suggests there exists an optimal prior strength value that balances bias reduction and computational cost, beyond which further increases in prior strength yield diminishing returns in bias improvement while incurring higher computational costs.

3.6.6 Summary / Recommendation.

Further analysis is needed to precisely identify the optimal prior strength value that minimizes bias variability while controlling computational cost.

3.7 Experiment 7: Prior Strength Derivation

Objective. To use the results from the experiment in Section 3.6 to derive a recommended prior strength value for BIOD based on desired bias characteristics. Based on the monotonic, sigmoidal behavior observed, we selected the hill equation as a suitable model.

3.7.1 Setup.

For the curve fits below, we will use the following hill curve function [2]:

$$y = y_{min} + (y_{max} - y_{min}) \frac{x^n}{k^n + x^n}$$

where:

- y : Normalized value at a given p_s .
- x : Prior strength p_s .
- y_{min} : Normalized value at $p_s = 1$.
- y_{max} : Normalized value at $p_s = 10^{12}$.




Figure 9: The compute / accuracy tradeoff in selecting p_s . Compute used is in the fraction of $k_{initial}$ actually performed in simulation, Bias is in the standard deviation of bias observed. Markers indicate recommended values for p_s .

Balance Factor	p_s	Compute	Bias	Δ Compute	Δ Bias	p_s Error
0.44	$7.08 \cdot 10^6$	0.83	$3.49 \cdot 10^{-2}$			$3.83 \cdot 10^4$
0.47	$4.71 \cdot 10^6$	0.8	$4.38 \cdot 10^{-2}$	$3.72 \cdot 10^{-2}$	$8.85 \cdot 10^{-3}$	$2.70 \cdot 10^4$
0.51	$3.35 \cdot 10^6$	0.76	$5.28 \cdot 10^{-2}$	$3.56 \cdot 10^{-2}$	$9.02 \cdot 10^{-3}$	$1.98 \cdot 10^4$
0.54	$2.48 \cdot 10^6$	0.73	$6.2 \cdot 10^{-2}$	$3.4 \cdot 10^{-2}$	$9.22 \cdot 10^{-3}$	$1.46 \cdot 10^4$
0.58	$1.9 \cdot 10^6$	0.69	$7.14 \cdot 10^{-2}$	$3.25 \cdot 10^{-2}$	$9.42 \cdot 10^{-3}$	$1.64 \cdot 10^4$
0.62	$1.49 \cdot 10^6$	0.66	$8.11 \cdot 10^{-2}$	$3.11 \cdot 10^{-2}$	$9.64 \cdot 10^{-3}$	$1.25 \cdot 10^4$
0.67	$1.19 \cdot 10^6$	0.63	$9.09 \cdot 10^{-2}$	$2.97 \cdot 10^{-2}$	$9.86 \cdot 10^{-3}$	$9.80 \cdot 10^3$
0.71	$9.7 \cdot 10^5$	0.61	0.1	$2.84 \cdot 10^{-2}$	$1.01 \cdot 10^{-2}$	$7.80 \cdot 10^3$
0.77	$7.97 \cdot 10^5$	0.58	0.11	$2.71 \cdot 10^{-2}$	$1.03 \cdot 10^{-2}$	$6.31 \cdot 10^3$
0.82	$6.62 \cdot 10^5$	0.55	0.12	$2.59 \cdot 10^{-2}$	$1.06 \cdot 10^{-2}$	$5.19 \cdot 10^3$

Figure 10: Table of optimal prior strength values for various balance factors β weighing bias reduction vs computational cost, based on hill curve fits to the normalized standard deviation of bias and normalized fraction of worst-case compute performed data.

- k : Half-maximal effective concentration, representing the p_s value at which saturation is at 50%.
- n : Hill coefficient, describing the steepness of the curve.
- Normalize the standard deviation of bias results to be a fraction of the maximum observed standard deviation across all prior strength values tested.
- Fit the hill curve function above to the normalized standard deviation of bias data

using non-linear least squares optimization to find the best-fit.

- Fit the hill curve function above to the fraction of worst-case compute performed data using non-linear least squares optimization to find the best-fit.
- Calculate the r^2 value for each fit to assess goodness of fit.
- Introduce a balance factor β and a Lagrangian objective function [1] to weigh the importance of bias reduction vs computational cost:

$$\text{objective}_{\text{score}}(p_s) = \text{bias}_{\text{std}}(p_s) - \beta \cdot \text{compute}_{\text{frac}}(p_s)$$

- For a range of 200 β values in the logarithmic space from 10^{-5} to 10^1 , find the prior strength p_s that minimizes the objective function.
- Only consider solutions where p_s is in the observed range of 10^0 to 10^{12} .
- For each optimal p_s found, find the error on p_s using the delta method:

$$\text{error}(p_s) = \frac{\sigma_{\text{hill}}^2_{\text{bias}} + (\beta \cdot \sigma_{\text{hill}}^2_{\text{compute}})}{\frac{\partial \text{objective}(p_s, \beta)}{\partial p_s}}$$

- Build a table of optimal p_s values for each β with the following columns:
 - (1) β : Balance factor.
 - (2) p_s : Optimal prior strength for the given β .
 - (3) Normalized fraction of worst-case compute performed at optimal p_s .
 - (4) Normalized standard deviation of bias at optimal p_s .
 - (5) Delta between current and previous row's normalized fraction of worst-case compute performed.
 - (6) Delta between current and previous row's normalized standard deviation of bias.
 - (7) The error on p_s

3.7.2 Definition of Metrics.

- p_s : Prior strength parameter provided to BIOD.
- Normalized standard deviation of bias: Standard deviation of bias normalized to the maximum observed standard deviation across all prior strength values tested.
- Normalized fraction of worst-case compute performed: Fraction of worst-case compute performed normalized to the maximum observed fraction across all prior strength values tested.

- r^2 : Coefficient of determination for the hill curve fits, indicating goodness of fit.
- β : Balance factor weighing bias reduction vs computational cost.
- n, k : Hill curve parameters as defined above.
- $p_{s,optimal}$ Recommendation prior strength for a given β .
- $\text{error}(p_{s,optimal})$: Estimated error on the recommended prior strength.

3.7.3 Hypothesis.

Based on the monotonic, sigmoidal behavior and apparent saturation effects observed in Section 3.6, we expect the hill equation to provide a good fit to both the normalized standard deviation of bias and normalized fraction of worst-case compute performed data, with fair r^2 values indicating strong goodness of fit.

We also expect to find a fairly wide range of optimal prior strength p_s values corresponding to a flat region in bias, with the error on p_s estimates being moderate, contributing to a wider recommendation range.

3.7.4 Results.

The hill curve fits yielded high r^2 values:

- Normalized standard deviation of bias fit: $r^2 = 0.9989$
- Normalized fraction of worst-case compute performed fit: $r^2 = 0.9991$

The table of optimal prior strength p_s values for various balance factors β revealed a range of p_s values corresponding to a flat region in bias where further increases in prior strength yielded minimal bias improvement while incurring higher computational costs.

It also revealed small errors (most $< 0.5\%$) on the values of both fits, and by extension, on the recommended prior strength values.

The table shown in Figure 10 shows a value of $p_s = 3.35 \cdot 10^6 \pm 1.98 \cdot 10^4$ provides a good balance between bias reduction and computational cost, yielding a normalized standard deviation of bias of 0.052 and a normalized fraction of worst-case compute performed of 0.76.

3.7.5 Analysis / Interpretation.

Both r^2 exceeding 0.998 confirm the hill equation is an excellent model for both the normalized standard deviation of bias and normalized fraction of worst-case compute performed data. The tight errors on the fits further reinforce the reliability of the model.

The table of optimal prior strength p_s values indicates a practical range for selecting p_s where the change in bias reduction becomes negligible compared to the increase in computational cost, allowing users to make informed decisions based on their specific bias tolerance and compute resource constraints.

The recommended value provides an option where increasing p_s to the maximum reasonable value yields 5% improvement in bias standard deviation at the cost of 24% increase in computational cost.

The $\beta = 0.51$ finding at the suggested p_s corresponds to a value determining the relative importance of bias reduction vs computational cost, that compute here is roughly half as important as bias reduction.

3.7.6 Summary / Recommendation.

We recommend a prior strength value of $p_s = 3.35 \cdot 10^6 \pm 1.98 \cdot 10^4$ for BIOD, balancing bias reduction and computational cost effectively for most applications.

The small errors on both the curve fits and recommended prior strength values provide confidence in the reliability of this recommendation.

4 Results

We evaluated the BIOD algorithm across a range of scenarios to assess its performance in terms of confidence accuracy, computational efficiency, and robustness to parameter mis-specification.

In the experiments described in Sections 3.1 and 3.2, we explored the effectiveness of BIOD’s initial failure probability estimation compared to simulated ground truth values. Figures 2 and 3 illustrate that BIOD’s estimates closely align with the actual failure probabilities, while remaining conservative, thus validating the algorithm’s reliability in practical applications.

In the subsequent experiment in Section 3.3, we tested the validity of the confidence interval BIOD uses to determine when to stop iterating. As shown in Figure 4, the lower-bound confidence intervals produced by BIOD consistently encompass the true failure probabilities across various dataset sizes and point-wise failure probabilities, demonstrating the robustness of the stopping criterion.

In addition, contrary to our hypothesis, the badly specified point-wise p values (blue and red curves) performed comparably to the well-specified values (green curves), indicating that BIOD is resilient to inaccuracies in initial parameter settings.

Next, in Section 3.4, we explored the effect of the damping constant on BIOD’s stopping condition via the floor on the number of iterations as set size increases. Figure 5 shows that the damping constant has a linear monotonic effect on the number of iterations performed, allowing for tuning based on available computational resources.

Finally, in Sections 3.5 through 3.7, we investigated the impact of prior strength on BIOD’s performance. Figure 9 shows that increasing prior strength leads to an improvement in bias of the failure probability estimates, albeit with diminishing returns, at the cost of increased computational effort following the hill curve for both metrics.

These results collectively demonstrate that BIOD is an effective and adaptable algorithm for iterative outlier detection, capable of providing reliable estimates while allowing for customization based on specific application needs.

5 Discussion

The results presented in Section 4 highlight several key strengths of the BIOD algorithm, as well as areas for further exploration.

One notable finding is BIOD’s robustness to mis-specification of the point-wise failure probabilities. Contrary to our initial hypothesis, the performance of BIOD remained consistent regardless of whether the point-wise p values were well-specified or deliberately mis-estimated. This resilience suggests that BIOD can be effectively applied in scenarios where precise prior knowledge of failure probabilities is unavailable, enhancing its practical utility in a wider range of applications than initially anticipated.

Two possible explanations for this robustness are the conservative nature of the initial failure probability estimates due to the damping fraction model, as well as the effect of the relatively strong prior strength used in the Bayesian updating process (Section 3.7). Both factors likely contribute to stabilizing the estimates against inaccuracies in the initial parameters.

Another finding was the lack of inherent bias in BIOD’s failure probability estimates across various scenarios (See Section 3.5). Across the universe of input parameters tested, BIOD consistently produced a mean bias near zero, with very little variation. This characteristic is particularly advantageous in applications where unbiased estimates are critical, as it ensures that the algorithm does not systematically overestimate or underestimate failure probabilities regardless of specified parameters.

The introduction of the damping fraction to temper the number of iterations as set size increases also proved effective in controlling computational cost while maintaining reliable estimates. The model represents a saturating monotonic relationship with C modulating the effective binding strength between outliers in the dataset. This allows users to tune the algorithm based on available computational resources without significantly compromising accuracy. It is noteworthy that p_s and C should be considered a pair - modifying C will likely change the optimal p_s value.

During the update procedure, instead of a pure Bayesian update after each iteration, we employed a bulk update approach, aggregating evidence from multiple iterations before updating the posterior distribution on failure. This has the advantage of eliminating the possibility of artificially lengthy tests in the case of an absence of failure without compromising the Bayesian nature of the update.

Finally, the use of a frequentist stopping criterion based on confidence intervals proved effective in ensuring that BIOD halts iterations once sufficient evidence has been gathered. The lower-bound confidence intervals consistently encompassed the true failure probabilities, demonstrating the reliability of this stopping condition. This approach balances the need for computational efficiency with the requirement for accurate estimates, making it a practical choice for real-world applications.

Future work could explore the extend of the independence assumption made in the BIOD model, particularly in datasets where outliers may exhibit correlations. Additionally, we investigated only the effects of normal and uniform outliers; investigating alternative prior distributions or adaptive prior strength mechanisms could further enhance the algorithm’s performance across diverse scenarios.

6 Conclusion

We have presented the Bayesian Iterative Outlier Detection (BIOD) algorithm, a method for adaptively determining the number of iterations required to achieve a desired confidence level in randomized testing. By combining Bayesian statistics with iterative refinement based on observed failure rates, BIOD improves upon traditional fixed-iteration approaches.

Empirical evaluation shows that BIOD is robust to poorly specified parameters and accurately models failure rates under independent, normally distributed errors. It balances computational cost with confidence in results, making it practical for resource-constrained settings such as CI pipelines. Future work could explore extending BIOD to correlated or non-normal error models, further broadening its applicability.

References

- [1] Hugh Everett. Generalized lagrange multiplier method for solving problems of optimum allocation of resources. *Oper. Res.*, 11(3):399–417, June 1963.
- [2] Rudolf Gesztelyi, Judit Zsuga, Balázs Varga, Bela Juhasz, and Arpad Tosaki. The hill equation and the origin of quantitative pharmacology. *Archive for History of Exact Sciences*, 66, 07 2012.

Appendix

Code

Rust Implementation

The repository for this paper contains a Rust implementation of the BIOD algorithm used for the experiments described. The code can be found at: https://github.com/rscarson/baysian_iterative_outlier_detection/

It is a ready-to-use implementation in the form of a crate named `biod`, which can be included in other Rust projects via Cargo.

Python Implementation

This code and a sample run can be found at: https://github.com/rscarson/baysian_iterative_outlier_detection/blob/master/biod.py

```
1 import math
2 import random
3 import sys
4 import time
5
6 from scipy.stats import beta
7
8 DAMPING_CONSTANT = 1.864
9 PRIOR_STRENGTH = 3.35e6
10
11 def p_fail_lower_bound(a, b, q):
12     """Calculates the lower bound of the (1-q) confidence interval for
13         a Beta(a, b) distribution."""
14     bdist = beta(a, b)
15     return bdist.ppf((1.0 - q) / 2.0)
16
17 def biod_k(q, p_fail):
18     """Calculates the required number of iterations k for BIOD given
19         confidence q and failure probability p_fail."""
20     return math.ceil(abs(math.log(1.0 - q) / math.log(1.0 - p_fail)))
21
22 def biod_run(data, f_transform, f_test, q, p, pass_ratio, timeout_s =
23             None, C = DAMPING_CONSTANT, p_s = PRIOR_STRENGTH):
24     """
25     Runs the BIOD algorithm on the given data.
26
27     Parameters:
28         - data: The dataset to run BIOD on.
```

```

26     - f_transform: Function to transform a set = fn(dataset, seed)
      -> transformed_dataset
27     - f_test: Function to test a set = fn(transformed_dataset) ->
      bool (pass/fail)
28     - q: Desired confidence level (e.g., 0.95 for 95% confidence)
29     - p: Pointwise failure probability estimate (should
      underestimate true failure probability)
30     - pass_ratio: Required pass ratio to consider the test
      successful (e.g., 0.99 for 99% pass rate)
31     - timeout: Optional timeout for the test function, in seconds (
      default: None)
32     - C: Damping constant (default: 1.864)
33     - p_s: Prior strength (default: 3.35e6)
34     """
35     start_time = time.monotonic()
36
37     # Calculate damping fraction and adjusted setwise failure
      probability
38     n = len(data)
39     damping_fraction = 1.0 / (1.0 + (C * n * p))
40     p_prime = p * damping_fraction
41     p_fail = 1.0 - (1.0 - p_prime)**n
42
43     # Calculate initial k
44     k = biod_k(q, p_fail)
45     k_initial = k
46
47     # Initialize priors
48     alpha = p_s * p_fail
49     beta = p_s * (1.0 - p_fail)
50
51     print(
52         f"Starting with initial pass count k = {k_initial}, "
53         f"confidence = {q * 100:.2f}%, "
54         f"required pass ratio = {pass_ratio * 100:.2f}%"
55     )
56
57     iterations = 0
58     passes = 0
59     unreported_passes = 0
60     last_failing_seed = 0
61     while iterations < k:
62         if iterations - passes > 0 and pass_ratio >= 1.0:
63             break
64
65         if timeout_s is not None and (time.monotonic() - start_time) >
            timeout_s:

```

```

66         print("\nTimeout reached, terminating BIOD run.")
67         break
68
69     # Perform transformation and test
70     seed = random.getrandbits(64)
71     transformed_data = f_transform(data, seed)
72     passed = f_test(transformed_data)
73
74     # Record result
75     iterations += 1
76     if passed:
77         passes += 1
78         unreported_passes += 1
79     elif passes < iterations:
80         # Failure detected, update priors and recalculate k
81         alpha += 1
82         beta += unreported_passes
83         unreported_passes = 0
84
85         last_failing_seed = seed
86         p_fail = alpha / (alpha + beta)
87         k = biod_k(q, p_fail)
88
89     #
90     # Final readout
91     #
92     failures = iterations - passes
93     p_fail_bound = p_fail_lower_bound(failures + 1, passes + 1, q)
94     print(
95         f"{q * 100:.2f}% confident that the true failure rate among
          datasets is at least {p_fail_bound * 100:.2f}%.\n"
96     )
97
98     ratio = (passes / iterations)
99     print(f"{failures}/{iterations} tests failed ({ratio * 100.0:.2f}%
100         pass)")
101
102     if iterations < k:
103         print(f"BIOD terminated early after {iterations} iterations -
104             not enough evidence to reach required confidence.")
105         return False
106
107     if ratio < pass_ratio:
108         print(f"BIOD failed to meet required pass ratio of {pass_ratio
109             * 100:.2f}%.")
110         return False

```

```
109     print("BIOD passed successfully.")
110     return True
```

Figure and Table Listings

1	Flowchart of the BIOD testing procedure.	6
2	Graph showing the conservative estimates of p_{fail} vs simulated true p_{fail} for various dataset sizes n	15
3	Graph showing the predicted number of iterations k to observed failure vs the simulated results	16
4	Graph showing the mean and standard deviation of lower-bound confidence intervals on the sample of iterations observed by BIOD from a simulation of 10,000 trials, compared to the simulated true p_{fail} for various dataset sizes n , and for various values of point-wise p . Green was deliberately underestimated, blue was exact, and red was deliberately overestimated. Markers indicate the number of iterations k performed by BIOD at each point.	18
5	Left: Graph of k_{floor} vs C values in the defined parameter space. Right: Table showing values of k_{floor} for $C_{\text{mean}} \pm C_{\sigma}$ for various confidence levels q	20
6	Graph demonstrating the lack of correlation between p_{fail} , pass_ratio , and bias in the base case of $p_s = 1$	22
7	Scatter plot showing the distribution of true p_{fail} values across all iterations of the experiment, demonstrating the coverage of the parameter space.	23
8	Comparison of bias behavior and convergence across p_s values	25
9	The compute / accuracy tradeoff in selecting p_s . Compute used is in the fraction of k_{initial} actually performed in simulation, Bias is in the standard deviation of bias observed. Markers indicate recommended values for p_s	28
10	Table of optimal prior strength values for various balance factors β weighing bias reduction vs computational cost, based on hill curve fits to the normalized standard deviation of bias and normalized fraction of worst-case compute performed data.	28